

Semantic Data Labelling & Type Safety

The Quest for sPandas

Riley Moher

What is semantic type safety?

CompanyA

Quarter	Profit	Operating Costs
Q1-2014	158 000 000	35 000 000
Q2-2014	200 000 000	38 000 000
...

CompanyB

Quarter	Profit	Operating Costs
Q1-2014	16	2.6
Q2-2014	30	3.3
...

[Profit Comparison] = CompanyA.Profit - CompanyB.Profit

- How do we ensure this computation is correct and meaningful?
- What could cause misinterpretation?
 - Different profit currencies
 - Different profit scale
 - Different quarters
 - Different profit definitions

What does the solution look like?

CompanyA

Quarter	Profit	Operating Costs
Q1-2014	158 000 000	35 000 000
Q2-2014	200 000 000	38 000 000
...

Company A
Unit
Definitions:

- Profit is net profit

Company B
Unit
Definitions:

- Profit is gross profit

CompanyB

Quarter	Profit	Operating Costs
Q1-2014	16	2.6
Q2-2014	30	3.3
...

sPandas Ontology Layer

IN:

```
CompanyA = SemanticDataFrame(data=CompanyA, units = CompanyAUnits)
```

```
CompanyB = SemanticDataFrame(data=CompanyB, units = CompanyBUnits)
```

```
[Profit Comparison] = CompanyA.Profit - CompanyB.Profit
```

OUT:

WARNING: Semantic dtype mismatch: 'profit type: expected {net} got {gross}'

How is type safety enforced?

- sPandas will contain a semantic units representation that underlies every semantic dataframe. These contain axioms which describe the different concepts that can be represented and how they are represented.
- When performing an operation requiring type-safety, types must be verified as semantically compatible:
 - a. First and foremost, they must be representing the same concept
 - b. They must be measured with the same scale and unit, or automatically converted to make this true (according to conversion semantics defined within the unit definitions)
 - c. They must have the same stratification(s), or be automatically converted to make this true (according to stratification semantics defined within the unit definitions)

Logic of Type-safety

- Components of units
 - Scale
 - Dimension (what are they actually measuring?)
 - Measure (what is the measure of the unit)
 - Has a numerical value
 - Domain-specific properties
 - Can be comparable or not
- Comparability
 - Properties of a given unit or dimension may be comparable or not
 - Non-comparable property means automatic conversion between distinct values of that property is not possible.
 - Some properties may only be comparable with additional information

Logic of Type-safety

- Using previous basic example of comparing profits of two different companies

CompanyA

- Profit(A)
- Profitkind(A,Net)
- hasScale(A, 1.0)
- hasCurrency(A,USD)

CompanyB

- Profit(B)
- Profitkind(B,Gross)
- hasScale(B, 10 000 000)
- hasCurrency(B,USD)

Unit Axioms

- hasMeasure(Profit, Money)
- NoncomparableDataProperty(ProfitKind)
- Dimension(Profit)

- We can convert the scale, and the currencies are comparable, but the different kinds of profits will cause a semantic error

Classifying sPandas Operations

- Different operations of sPandas can be classified in 3 ways
 - Operations enforcing type safety
 - Operations producing new types
 - Type neutral operations

Type-Safe Enforced Operations

- Addition/Subtraction
 - `dfA['col1'] +/- dfB['col2']` -> verify concept being represented is the same, automatically change scale/unit if conversion definition exists, otherwise raise semantic type mismatch error
- Filtering
 - `df[df['price'] >= sPandasUnit{value: 10M, scale: 1, curr: USD, agg: {mean across 7 days}}]` -> check for equivalence of each price row according to the semantics of equivalence for that unit, performing any automatic conversions if applicable.
- Concat, append, join, merge
 - `df1.join(df2), df1.append(df2)` -> any columns from one dataframe being added to another must be verified for compatability

Type-Changing Operations

- Groupby
 - Will attempt to convert units of columns other than the one being grouped-by according to groupby aggregation method invoked
- Statistic operations: mean, count, max, min, median, std
 - Concept being represented is the same, but stratification is put in place according to the statistic
- Apply
 - Any functions called with `df.apply(func)` must have a static-typed return
- Importing
 - Importing will automatically apply typing to certain variables (binary, categorical), will convert other units to basic dimensionless unit type
- Multiplication/Division
 - New units defined through geometric combination of multiple units, ex: $\text{cm}^2 == [\text{cm}][\text{cm}]$, $\text{\$/person (annual)}$,

Type-Neutral Operations

- Sorting operations
 - Sorting changes order, no type awareness necessary
- Exporting
 - Exporting will behave as in pandas, could possibly export some other metadata file for units?

Inference

- One of the advantages of using an ontology-based approach is the ability to leverage an automated reasoner
 - Can check if a given unit class subsumes another
 - Can also use multiple inheritance and multiple instantiation
 - Can also find units that are equivalent, if subsumption check is true in both directions (automatically done with Hermit reasoner)

Unit Types

- Units of Measure
 - Include familiar scientific units of measurement for length, volume, area, etc.
 - Also includes measurable concepts like money-based concepts, profit, GDP, etc.
- Time dimension
 - Quantification of time-based validity
 - Can be combined with units of measure
 - Could be instantaneously defined (timestamped measurement) or defined over a time-period
- Statistical dimension
 - Representation of aggregational operation(s)
 - Can be combined with units of measure
- Geospatial dimension
 - Similar to time and statistical, modifier for measure unit
 - Defined over some geographical space

Unit Types

- Categorical unit
 - Basic unit to define categorical units
 - Semantics of these units will be harder to generalize, could provide some basic framework to be extended on a per-case basis

Avoiding Consequences

- Through real failure analysis in 3 domains, we can show how a system like sPandas can mitigate real consequences

Predicting Congestive Heart Failure with 100% Accuracy?

- In [a paper published January 2020](#), Porumb et al present a CNN model that accurately identifies congestive heart failure (CHF) on the basis of one raw electrocardiogram (ECG) heartbeat only
- The 100% accuracy boasted is misleading, heterogeneous data sources and failure to properly validate caused a big problem
 - 100% accuracy is on the training dataset, 97.8% accuracy on testing set
 - Only 33 subjects as datasources
 - **Positive and negative samples came from different ECG sources, with different sampling frequencies**
 - Data sources were all from lead-1 ECG (single lead)
- Having data source information encoded within the data could have prevented this inaccurate analysis
 - When analyzing the model's predictions, semantic information about the data source could have revealed biases
 - Semantic data provenance would have revealed the artificial downsampling difference in the positive vs negative classes

Use Case 1: Python for Finance - Time Series Analysis

- Data sources are TSLA financial information and US macroeconomic data
- Utilizes many of pandas' timeseries functions
 - Resample
 - Shift
 - Rolling
- Uses statistical models for time series modeling
 - ETS
 - EWMA
 - ARIMA
- Will be useful to demonstrate time series representations
 - Aggregations over periods of time
 - Timeseries data with an implied causality

Use Case 1: Python for Finance - Time Series Analysis

What's the challenge?

- All timeseries information is contained within the index
- No data provenance across timeseries operations like shifts, reindexing, etc
 - Units transformed through these operations will have some assumptions in the way they're aggregated, these assumptions are lost through these transformations

How can sPandas address it?

- Integrating timeseries information in the units
 - Ex: `print(stock['price'].dtype)` → sPandas Unit {scale: 1, curr: USD, agg: {mean across 7 days}}
- Automatic data provenance through operations

Use Case 2: Cervical Cancer Risk Classification

- Data contains many different categorical variables
- No intuitive units for this data, physical quantities would not be useful
- Good example of filling missing categorical data
 - What are the semantics for values filled by taking median, etc?
- Contains many ML data formatting and processing tasks
 - Can be used to evaluate sPandas

Use Case 2: Cervical Cancer Risk Classification

What's the challenge?

- Categorical values may have different names for same concepts - large vocabulary in medical field
- Filling missing values contains assumptions that are not transparent and are not easily traceable
- Normalization of values reduces information in data, provenance not maintained

How can sPandas address it?

- Automatic conversion of categorical variables to a categorical unit type, which can be later refined to some specific concept
- Ontological definition of medical terms from a medical domain ontology
- Integrated data provenance

Use Case 3: Identifying Invalid GPS Taxi Data

- Combines geospatial with timeseries
 - A good opportunity to integrate both representations
- In this specific use case, pandas should make the task of identifying invalid trip data much simpler
 - Semantics of geospatial units will have axioms that are analogous to the functions being manually implemented in this use case
 - Reduces work required in geospatial data cleaning

Use Case 3: Identifying Invalid GPS Taxi Data

What's the challenge?

- Same challenges as in use case 1 with regards to the timeseries element
- Identifying outliers in geospatial data requires considerable manual work
 - Manually defining a filtering function
 - Requires a high degree of exploratory data analysis that could be reduced with some base-level geospatial semantics

How can sPandas address it?

- With embedded geospatial semantics, tasks like filtering out invalid GPS data will be much faster
 - Semantics of distance, travelling, etc.
 - Axioms will be able to error-check data like this with much less effort

Implications for Reproducible Research

- Automatic data provenance, especially as an extension of a widely-used tool in data science, makes results clearer to follow, mistakes easier to spot.